

1 Simulating GPS-spoofing of consumer UAVs as a
2 side-channel remote control method

3 Botan Sinayic
bs477@kent.ac.uk



School of Computing
University of Kent

Word Count: 6,000

4 April 2, 2025

5 Contents

6	1	Introduction	2
7	2	Background & Related Work	3
8	2.1	Prior research	6
9	3	Design	7
10	3.1	Simulation setup	7
11	3.2	Why SITL?	7
12	3.3	Environment components	8
13	3.4	Simulation environment	9
14	4	Attack Implementation	10
15	4.1	EKF and Offset	12
16	4.2	First implementation	13
17	4.3	Second implementation	13
18	4.4	Controller	16
19	5	Results & Analysis	18
20	6	Discussion	19
21	6.1	Countermeasures	19
22	6.2	Technical countermeasures and defences	19
23	6.3	Operational and policy strategies	20
24	6.4	Future works	21
25	7	Conclusion	21
26	7.1	Adversarial threats and abuse scenarios	21
27	7.2	Validation Framework for Enhancements	23

29 With the rapid growth of UAV technology, applications in civil domains,
30 cargo transport, and other fields have expanded. As we find new applica-
31 tions for these technologies, new cybersecurity threats emerge. It is essential
32 to protect such technology from cyber attacks, to ensure safe, secure, and
33 reliable usage for all.

34 This framework can serve as a dual use simulation tool that can allow
35 drone manufacturers to evaluate the resilience of Commercial off the Shelf
36 (COTS) positioning components against GPS spoofing attacks by evaluating
37 known, modelled hardware and software. However, for the development of
38 new models, such as adjustments to the Extended Kalman Filter for fine tun-
39 ing will require real world validation. This simulation may not fully capture
40 potential false positives compared to real world as it does not include en-
41 vironmental conditions such as weather. Ultimately, the simulation demon-
42 strates how differential GPS spoofing combined with the use of controller
43 APIs can allow an attacker to remote control a drone in some scenarios by
44 translating X-Y coordinates into spoofed differential GPS coordinates. Due
45 to legal, safety and ethical constraints, all experiments were conducted in a
46 safe, hardware free environment.

47 Although carrying out experiments on real hardware would yield more
48 accurate results, simulation offers an alternative which is reflective of real
49 hardware conditions, legally compliant, and safe approach. Our results sug-
50 gest that a stealthy GPS spoofing attack against Arducopter can enable an
51 attacker to gain full remote control of the drone without the need to hijack
52 the control channel, which is typically encoded and encrypted, making it a
53 poor target for real-time attacks. Instead, the attack will leverage manipu-
54 lated GPS signals to mislead the drone about its true position.

1 Introduction

A GPS spoofing attack involves an attacker crafting a malicious GPS signal that tries to replicate a legitimate signal to deceive the flight controller to miscalculate its location. This is because GPS signals are not encrypted, due to the difficulty of coordinating key exchanges between satellites and ground receivers. Receivers have insufficient power to respond to signals received from satellites in a GPS, GLONASS or other positioning constellations. GPS receivers lock onto any GPS signal that is sufficiently powerful, well-crafted and typically require a minimum of 8 confirmed satellites, with 13 being optimal to ensure accurate positioning and reduce errors. Arducopter, however, required a minimum of 6 satellites. In this project, the Software-in-the-Loop (SITL) simulation framework will be employed to evaluate the resilience of UAV flight controllers against malicious GPS signals. The simulation uses open-source Arducopter module part of the Ardupilot suite as the test drone and will replicate real-world sensor and flight behaviours in a controlled, hardware free environment.

At the core of our methodology is the exploitation of the Extended Kalman Filter (EKF), which integrates sensor fusion techniques to detect discrepancies by observing changes in GPS, magnetometer, an Inertial Measurement Unit (IMU) data over time. This fusion enables the EKF to determine whether the drone's position has shifted without corresponding instructions from the flight controller, a potential indication of a dangerous fault. By combining X-Y coordinates with incremental offsets to the legitimate GPS data, the framework demonstrates how the EKF's tolerance threshold can be exploited without triggering its safety measures.

GPS spoofing represents a serious security and safety threat to consumer UAVS and any technology that relies on GPS data for positioning. For instance, a drone on an autonomous mission can be hijacked or misdirected for a malicious purpose. Detailed abuse scenarios of such attacks will be discussed in the conclusion section.

A contribution of this project is the integration of a PS4 controller to

86 provide a user-friendly interface for employing GPS spoofing to manually
87 control and idling drone. This contribution will yield insights into how an
88 adversary might combine GPS spoofing with real-time controller inputs to
89 remotely hijack a UAV. This framework offers a legally compliant and safe
90 method of testing UAV security measures and contributes to the development
91 of powerful and effective countermeasures against cyber threats.

92 2 Background & Related Work

93 Prior to undertaking the simulation, we first explored the possibility of using
94 the USRP B210. A detailed explanation of this attempt can be found in
95 Appendix C.

96 The Global Positioning System (GPS) is a critical navigation system that
97 relies on GPS signals received from satellites on the L band. Civil GPS signals
98 are not encrypted, largely because encrypting them is a hard design problem
99 given the extensive infrastructure required for reliable ground-to-space com-
100 munication. As a result, it is not feasible to encrypt this data. Moreover, the
101 integration of sensor fusion such as the EKF in conjunction with methods like
102 Google’s positioning infrastructure and cellular triangulation makes encryp-
103 tion redundant. Drones do not benefit from 5G connectivity yet, position
104 verification must rely on EKF’s ability to combine position related sensor
105 data to identify anomalies.

106 GPS operators work by receiving signals from multiple satellites in medium
107 Earth orbit, each satellite being approximately 20,000 km away. By measur-
108 ing the time delay between when the signal was sent and when it was received,
109 the GPS receiver computes the pseudo-distance, or the approximate distance
110 to each satellite. With signals from at least four satellites, the receiver uses
111 trilateration to determine its three-dimensional position.

112 Extended Kalman Filter in navigation systems is a crucial component in
113 all modern flight controllers like Arducopter. It is an algorithm that estimates
114 the drone’s position, velocity, and altitude. The EKF not only tracks where

the drone is through this continuous estimation process but also identifies any sensor inconsistencies. Comparing the predicted value and the latest (non-IMU) value every 100 ms, the EKF is able to calculate the innovation value and the variance value. The variance value reflects how much trust the EKF has in its current estimate. If the innovation values are positive and the variance value is 1 or higher, a fail-safe is triggered. A sudden GPS jump without corresponding movement detected by the IMU indicates a potential sensor fault, in such cases the flight controller triggers a fail-safe.

Arducopter requires a minimum of 6 confirmed satellites to obtain an accurate GPS fix. This redundancy is required as it compensates for signal loss, interference, and environmental factors. For example, multipath effect occurs when two radio signals reach the receiver via multiple paths, and this causes interference and signal degradation, impacting the reliability of positioning. Even though, environmental factors do not exist in the simulation environment, the simulation tries to mimic this type of interference through configurable sensor noise and error injection models. For example, our simulation setup includes parameters such as `SIM_ACC_RND` to add noise to the accelerometer, `SIM_WIND_SPD` and `SIM_WIND_DIR` to simulate wind effects. These types of models deliberately introduce noise, random errors and distortions to replicate real world conditions like signal delays, sudden drops, and inaccuracies. These models and approach enable manufacturers to test and refine the autopilot's filtering and fault detecting algorithms like the EKF.

For instance, if wind causes the copter to drift, the flight controller will rely on the EKF to fuse GPS and other sensor data to update the vehicle's position estimate so that path corrections can be made. An attacker can advantage of this behaviour by injecting fake incremental drifts in the spoofed GPS packet that mimics natural sensor noise and remains below the EKF's threshold, by doing this the flight controller will interpret this change as normal drifts and will attempt to correct its course based on its last known position. By continuously transmitting these spoofed GPS packets the at-

146 tacker will be able to steer the copter in their desired location. It is important
147 to note that the EKF itself does not directly control complex/correctional
148 behaviours, it provides an estimate of position based on fusion of sensor data
149 that the flight controller uses to make the course corrections. If the attacker
150 exceeds this threshold, the EKF will trigger a fail-safe and land. So the suc-
151 cess of a GPS spoofing attack depends on keeping under this threshold set
152 by the EKF

153 Most modern commercial drones on the market rely on sensor fusion al-
154 gorithms like the EKF, however, while the EKF is widely implemented, its
155 configurations like the threshold for errors is highly dependant on the spe-
156 cific sensor suite and hardware characteristic of the drone. For instance, a
157 larger more powerful drone will introduce more vibrational noise, therefore,
158 the IMU sensors would need have more enhanced shielding and a higher
159 polling rate for the EKF to effectively filter out the noise while still pro-
160 ducing a precise positional estimate. Consequently, the drone manufacturers
161 tend to keep these specific configurations and behaviours of their sensor fu-
162 sion algorithms confidential and also put in countermeasures to make reverse
163 engineering more difficult for attackers. As a good understanding of these
164 mechanisms will equip the attacker with the necessary knowledge to develop
165 bypass methods.

166 The inertial measurement units, such as the accelerometer and gyroscope
167 measure the linear acceleration and angular velocity. It estimates short-
168 term changes in position and altitude. The challenge with this is that the
169 accelerometer suffers from drift, over time the raw data leads to errors. The
170 magnetometer measures the orientation of the drone relative to the earth
171 magnetic field and this is susceptible to electromagnetic interference. These
172 are the reasons why a tolerance threshold is implemented as these drifts
173 occur naturally and the flight controller needs to consider them as they can
174 accumulate to greater errors.

175 GNSS systems are associated with specific geopolitical regions, they are
176 controller by various political entities. For example, the United States op-

erate GPS, Russia controls GLONASS, and the European Union manages Galileo. These entities have the ability to disable and modify signals, this drives the development and development of independent GNSS satellite systems.

2.1 Prior research

In the well-known UT Austin Humphreys et al. (2012) experiment, researchers successfully managed to capture a drone using GPS spoofing using real hardware. They also managed to make the UAV have a large number of unrecoverable navigation errors that caused the UAV to crash. Gupta et al. (2023) adopted the simulation-based approach, by leveraging the pre-built module called `fake_gps`. The keys features of this module include:

- **MAVLink integration:** This module uses the MAVLink protocol to communicate GPS data, which allows for precise injection of offsets into the telemetry stream. This capability facilitates a controller simulated of GPS spoofing attack. Even though big manufacturers tend to opt out of propriety software like DJI's occusync, MAVLink integration is still common in hobbyist drones and those that leverage open source autopilot systems like Ardupilot and PX4.
- **Offset injection:** Now that the drone communicates via MAVLink, the module now acts as a bridge to the GPS telemetry stream, this enables us to simulate various GPS spoofing scenarios in a controlled manner.
- **Ease of use:** Because it is already built in, it wont require extensive custom development, facilitating a quick setup and experimentation.

In the project Noh et al. (2019), the authors took the approach of directly modifying the GPS emulator part of the HAL abstraction layer of the SITL. This approach allowed the researchers to interact with the underlying simulation at a low level, this ensures that any modification effect the system

205 in a manner that closely mimics how changes would manifest in real hard-
206 ware. By making modifications to the HAL abstraction layer, the entire data
207 flow between the simulated sensors and the control algorithm is influenced,
208 enhancing realism.

209 Another approach includes using the pre-built parameters/commands,
210 `SIM_GPS1_GLTCH_X`, `SIM_GPS1_GLTCH_Y` which applies a apply offsets to the
211 latitude and longitude. This was built by Arducopter developers to test the
212 effects of wind and other types of drifts. We will be taking this approach in
213 our first implementation. The key features of taking this approach includes:

- 214 • **Ease of integration:** As these parameters are already built into Ar-
215 ducopter framework, they can be called without the need for extensive
216 code modifications or additional custom modules.
- 217 • **Precise offset control:** These parameters allow the injection of exact
218 offset values to the latitude and longitude.
- 219 • **Reproducibility:** This approach allows experiments to be consistent
220 and repeatable, which is crucial for systematic testing and validation.

221 3 Design

222 This section will entail the design and execution of the simulation-based
223 experiment, describing the simulation setup, rationale for using Software-in-
224 the-Loop (SITL) environment and the attack implementation strategies.

225 3.1 Simulation setup

226 3.2 Why SITL?

227 To be able to conduct this experiment in a safe, ethical and legally compliant
228 manner, we employ the ArduPilot SITL (Software-In-The-Loop) simulator¹.
229 This framework can utilise the open source Arducopter module from the

¹<https://ardupilot.org/dev/docs/sitl-simulator-software-in-the-loop.html>

230 Ardupilot suite as the test drone, enabling us to replicate real sensors and
231 flight behaviours in a controlled, hardware free environment. This setup will
232 enable us to have precise control over the experiment conditions and a safe
233 way to inject malicious GPS data directly into the simulated flight controller,
234 without altering the drones actual simulated world position. The physical
235 location simulation will remain consistent, while the flight controller responds
236 to spoofed GPS readings as though they were genuine. This approach ensures
237 we can precisely manipulate the drone’s perceived location and evaluate its
238 behaviour and the effectiveness of the safety measures in place.

239 Using this framework comes with several advantages. It ensures legal
240 and safety compliance by eradicating any risk associated with the emission
241 of unauthorized signals on real hardware. It also provides flexibility to the
242 project, enabling us to make quick iterations and adjustments to the attack
243 parameters such as offsets and rate. Finally, SITL enables a high-fidelity
244 replication of flight dynamics and sensor behaviour, making it a powerful
245 proxy for real-world testing.

246 3.3 Environment components

247 • Arducopter ²

248 Arducopter will be utilised as the flight control firmware.

249 – Open-source nature: Arducopter is a widely used open source
250 module that enables us to include extensive modification and re-
251 alistic of UAV behaviour.

252 – Built-in EKF: More close to real world conditions as modern
253 drones utilise countermeasures like the Extended Kalman Filter
254 algorithm.

255 • QGroundControl³

256 We will use this as our ground controller.

²<https://ardupilot.org/copter/> Arducopter documentation

³<https://qgroundcontrol.com/>

- 257 – Easy quick setup
- 258 – Visualization tools: Allows us to view telemetry data, parameters
- 259 and sensor outputs in real time.
- 260 – Easy installation of flight logs
- 261 • Packages
- 262 We will use these packages for our spoofing script.
- 263 – PyGame⁴ - utilised to retrieve analogue inputs from the PS4 con-
- 264 troller.
- 265 – DroneKit⁵ - This will serve as the communication bridge between
- 266 the simulation and Arducopter's flight controller.

267 3.4 Simulation environment

268 For this project we will be using Ubuntu 22.04 virtual machine running on
269 an Apple M1 Pro. To launch the Arducopter simulation run,
270 `sim_vehicle.py -v ArduCopter --console --map --osd`
271 `sim_vehicle.py` is a startup script that configures the current code branch
272 of the SITL firmware, this includes the vehicles position, altitude, orienta-
273 tion, and velocity. This process defines both the simulation world space and
274 the entities within it. The world space is incorporated with global physic
275 parameters such as gravitational forces, wind conditions and air resistance,
276 these parameters effect all of the objects within the simulation. The copter
277 itself is an individual object that has specific state parameters and interacts
278 with the environment according to the simulated physic laws.

279 Once the simulation is running, the vehicles state is continuously updated
280 by the SITL physics simulator, it computes the current state of the copter us-
281 ing motor output values. This data is then retrieved by the `SIM_GPS_GPS::read()`
282 function and forwarded to the GPS manager `AP_GPS`, updating its GPS object

⁴<https://devdocs.io/pygame/>

⁵<https://dronekit.io/>

every 200 ms with fresh positional data. The EKF then retrieves this object and fuses it with IMU sensors to generate a state estimate which takes account for noise and sensor variances.

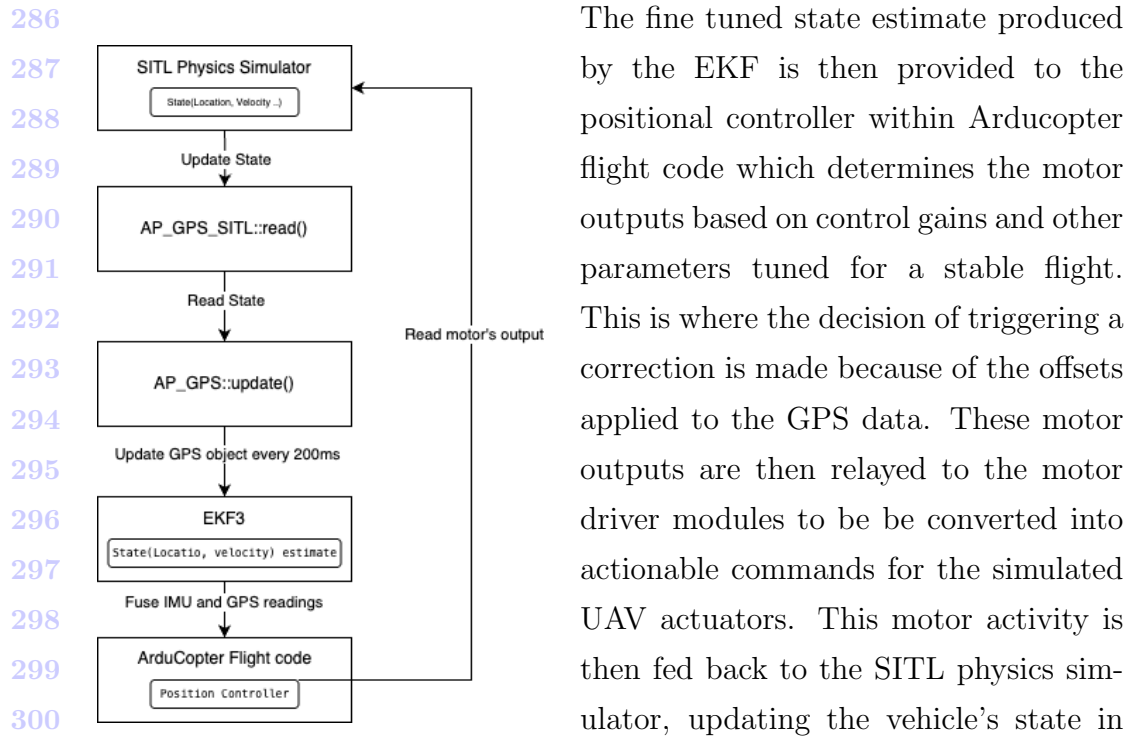


Figure 1: Execution loop of SITL environment

In the context of the GPS spoofing attacker, rather than altering the vehicles state provided by the physics simulator, we alter the updates made to the GPS object by the `AP_GPS` module. This approach attempts to simulate real-world conditions where spoofed sensor inputs must pass through validation checks, such as the Extended Kalman Filter sensor fusion process before being able to influence the flight controller.

4 Attack Implementation

Further detailed explanation of earlier experiments and an analysis of results can be found in Appendix A and B.

312 The primary objective of this attack is to manipulate Arducopter per-
313 ceived position by injection spoofed GPS signals in a way that a path cor-
314 rection is triggered by the flight controller causing it to steer in the attackers
315 desired location. All while maintaining stealth and within the tolerance limits
316 of the EKF to bypass safety mechanisms such as a fail-safe. We will attempt
317 to achieve our objectives through two different approaches, all essentially
318 attempting to inject spoofed GPS data to the GPS telemetry stream.

319 Early experiments which included attempts to spoof the GPS with arbi-
320 trarily large leaps resulting in immediate detection, this insight resulted in
321 further investigation into the role of the EKF and the process of fusing GPS
322 data with Inertial Measurement Unit sensor outputs. It became clear that
323 the EKF is designed to filter out sudden, unrealistic changes by comparing
324 different sensor outputs every 100 ms and triggering safety mechanisms in
325 case of accumulation of discrepancies which exceed the predefined threshold.
326 This led the attack implementation to adopt a strategy of gradual, incremen-
327 tal offsets that mimic natural sensor noise, thereby gaining stealth.

328 **First Implementation:** In the first implementation, we will be using
329 pre-built parameters called `SIM_GPS_GLITCH_X` and `SIM_GPS_GLITCH_Y`, each of
330 which adds a specific pre-determined offset to the GPS data being sent to
331 the flight controller. We will keep track of each offset applied and apply
332 the offset to the spoofed location, which will cause the effect of *ArduCopter*
333 chasing the spoofed perceived location. The script will be run while the UAV
334 is in *Auto* and *Guided* modes, both of which rely heavily on GPS.

335 **Second Implementation:** The second implementation involves using a
336 pre-built module within the Ardupilot suite called `fake_gps`. This will allow
337 us to inject spoofed GPS data via MAVLink communication. To force Ar-
338 ducopter to use MAVLink as its communication protocol we use the command
339 `param set GPS1_TYPE 14`. We will be modifying the script to integrate the
340 controller API and make minor adjustments.

341 Step by step instructions on how to reproduce these attacks have been
342 included in this project's repository README file, the link can be found in

the footnote⁶.

4.1 EKF and Offset

To carry out a successful GPS spoofing attack against consumer drones, we first need to have a good understanding of the fail-safe mechanism in place. The Extended Kalman Filter plays a crucial role in this, by monitoring and keeping track of discrepancies between the GPS data and IMU sensor outputs, it triggers a safety measure if it exceeds a pre defined threshold.

By fusing the IMU output with the GPS data, the EKF algorithm predicts velocity, position, and magnetic field. IMU sensors inherently have errors, which the EKF algorithm takes into account. For instance, the accelerometer has a constant bias error which causes a position error that accumulates over time. The gyroscope has angle random walk error, it will exhibit high-frequency white noise due to thermoelectric reactions. Due to these errors, EKF has to count for them. It does this by comparing the data every 100ms and a fail-safe will be triggered if the error count exceeds 10. If this threshold was not implemented, the EKF algorithm would constantly have false positives and trigger fail-safes. The accuracy of the sensor outputs has a large influence on the effectiveness of the EKF and different sensors on different drones exhibit varying levels of accuracy. Therefore, different drones have different degrees of errors and this needs to be taken into consideration for the most effective EKF.

An attacker can exploit this tolerance by injecting spoofed GPS data that fall in line with the threshold accepted by the EKF. This allows us to inject fake GPS data that seem normal to the EKF, thus, deceiving it to believe that it is at a certain position triggering a path correction in the attacker's desired location.

Prior research suggested that Arducopter changes flight mode on activation of a fail-safe, however, this was not the case for our first implementation and we will dive deeper on to the reasons why. The second implementation,

⁶<https://github.com/Botii/gpsSpoofingProject>

however, aligns with the findings from prior research and Arducopter did switch to Land mode on the trigger of a fail-safe.

4.2 First implementation

Preliminary experiments related to the second implementation, first implementation and the controller feature are detailed in the Appendix. Please refer to the Appendix for additional background and supporting information on these early experiments.

Results from the first implementation showed that the path-following algorithm detected a deviation, and the Extended Kalman Filter (EKF) switched sensor reliance from GPS to IMU. A lane-switch corrective action was observed seven times during the flight before reaching the waypoint. This is what we call a soft failure, when the EKF switches from relying on GPS to IMU sensors without detecting a critical error sufficient to trigger a full fail-safe. Only one sensor became unreliable, while others, such as the IMU, remained a reliable alternative for maintaining an accurate state estimate. Even though the EKF relied more on the IMU sensors, the drone carried on being spoofed. I believe this is a simulation bug, as the GPS data still had enough weight to alter the state estimate.

4.3 Second implementation




Figure 2: Large offset being detected by the EKF

In our first experiment, we applied very large offsets, which immediately caused the EKF to trigger a fail-safe. For example, applying a sudden offset

393 of 0.0001 was quickly recognized by the EKF, resulting in a mode switch to
 394 landing. As shown in Figure 3, the green line (square root of the velocity
 395 variance), blue line (innovation in velocity), and orange line (innovation in
 396 position) all exceed the EKF threshold of 0.8. This triggers a fail-safe, caus-
 397 ing the copter to switch to land mode and landing. In contrast, Figure 4
 398 illustrates a normal flight in guided mode without any spoofing, the values
 399 are much lower with a maximum of only 0.08 compared to a peak of 3.9
 400 during the spoofed flight in Figure 3.




Figure 3: Regular flight with no offsets

401 The EKF threshold can be adjusted via `FS_EKF_THRESH` parameter, the
 402 EKF's fail-safe action can also be changed via `FS_EKF_ACTION` parameter.
 403 Land fail-safe action may not always be the best option, an alternative can
 404 be `AltHold`.




Figure 4: Incremental offsets

405 This initial result prompted further investigation into reducing the mean
 406 of the EKF values to similar values seen in Figure 4. One approach is using
 407 incremental small offsets, with an incremental offset of 0.000001 we managed

to reduce the mean to 1.19 allowing the attacker to spoof the copter some distance before a trigger of a fail-safe. This gradual increase removed sudden spikes and maintained control of the copter until the variance eventually reached the threshold, at which point the EKF triggered a fail-safe as seen in Figure 5.




Figure 5: Applying offsets incrementally and in bursts

Further investigation led the approach to applying the offsets in bursts to remove the cumulative error. Through trial and error with various burst interval, we determined that burst intervals of 1.8 provided the best stealth performance. As shown in Figure 6, none of the three monitored EKF parameters exceeded the threshold at any point. This allowed us to achieve a mean value similar to Figure 4, the mean remained below 0 during the whole flight. This approach effectively hijacks the arducopter until battery depletion. However, the limitation is that this technique only works when modifying the latitude.

This is caused by the limitation of the magnetometer. Near the north and south poles, the Earth’s magnetic field lines become nearly vertical. It relies on detecting variations in the horizontal component of the magnetic field to detect movement. When the copter is travelling in the north-south direction, the horizontal variance is minimal because the field lines are almost parallel. The magnetometer is unable to accurately calculate its relative position or detect movements, making it less sensitive to spoofing in these directions.

The pre-defined magnetic field variance threshold is set to 50 by default ⁷.

⁷<https://ardupilot.org/copter/docs/parameters-Copter-stable-V4.5.7.html>

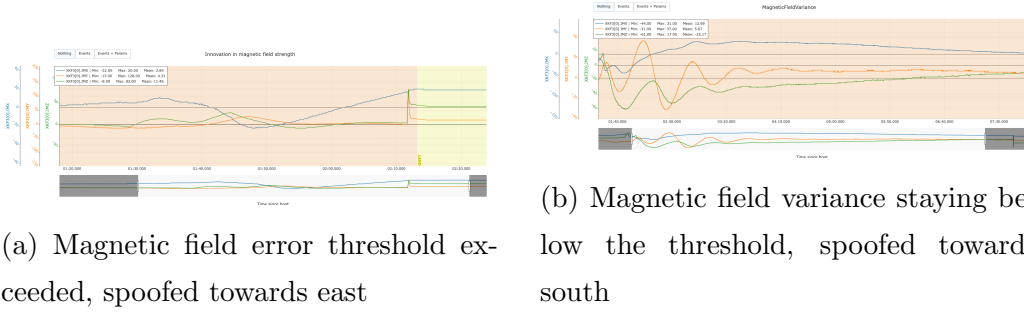


Figure 6: A side by side comparison of magnetic field variance varying based on direction of spoof

When this threshold is exceeded, a fail-safe is triggered. As shown in Figure 7a, the magnetic field innovations in both the X-Y and Z axis exceeds the threshold of 50, triggering a fail-safe. In contrast, when spoofing the copter in the southward direction, we can see from Figure 7b the magnetic field innovations never exceeds the threshold. This discrepancy occurs because the magnetometer is unable to accurately determine its relative position near the poles due to the nearly vertical orientation of the magnetic field. The lack of significant horizontal variation means that the EKF does not register a spoofing event in these conditions.

4.4 Controller




Figure 7: Flowchart of the GPS spoofing script with controller-based offset

Due to the success achieved from the second experiment in the appendix without the controller, I implemented the controller integration to determine if we can control the copter within a radius of 97.98 m without detection by the EKF.

I created a new function called `get_controller_input` which is responsible for retrieving the X and Y axis from the controller API. I applied a dead-zone to prevent unwanted drift from joystick noise, this is necessary as




Figure 8: Latitude and longitude changes over time

447 it causes the EKF to trigger a compass error. I've also reduced the applied
 448 offset scale from 0.000001° to 0.0000005° . The script continuously adds lat-
 449 itude and longitude offsets based on the real-time controller input, allowing
 450 dynamic control of the spoofed location while maintaining stealth.




Figure 9: Sudden change in direction on the joystick causing fail-safe




Figure 10: EKF values exceeding the threshold

451 I ran this script several times with varying joystick input. In the first
 452 experiment, sudden changes in the joystick input caused the EKF to detect
 453 the abrupt offsets. As shown in Figure 9 (red box), a sudden increase in
 454 longitude triggered a switch to land mode, a fail-safe response by the EKF.
 455 Looking at the EKF values for the same flight, which can be seen in Figure
 456 10, the sudden increase in longitude caused the velocity variance to exceed
 457 the threshold of 0.8.

458 This occurs due to the controller-induced lag, during which the offsets
 459 accumulate, when the GPS data is finally processed, the accumulated offset
 460 is applied abruptly rather than as gradual changes.

461 However, with smooth joystick movement, gradual and consistent offsets
 462 were applied to latitude and longitude, as shown in Figure 8 enabled us to
 463 successfully take full control of Arducopter via GPS spoofing. From the take-




Figure 11: EKF values staying under the threshold



Figure 12: Achieving 3,189 m distance with GPS spoofing attack

off point until the UAV crashed, we managed to reach a distance of 3,180.4 m or 3 km. As shown in Figure 11, the EKF values stay below the threshold of 0.8 during the whole period of the attack.

5 Results & Analysis

In the first implementation, we were able to partially hijack arducopter, eventually triggering an EKF lane switch when a specific distance was reached. However, even though multiple lane switched would occur, we were still able to spoof Arducopter. I believe that this is a simulation bug, as the GPS data still held enough weight for the EKF to produce spoofed state estimates.

In the second implementation, we successfully spoofed Arducopter in the north and south direction until battery depletion using the burst technique. However, any other direction caused the EKF to trigger a fail-safe once a threshold was reached.

The second implementation with the controller feature (not using the

burst technique), allowed the UAV to be fully hijacked without triggering any safety mechanisms if the attacker maintains smooth joystick movement eliminating sudden spikes in EKF values.

6 Discussion

6.1 Countermeasures

In this section we will be exploring the different ways in which we can defend against GPS spoofing attacks. I will be considering both technical solutions which will be built into the drones and higher-level policies and practices.

6.2 Technical countermeasures and defences

The most obvious answer would be to utilise cryptography and encrypt all GPS signals. This is not common practice in consumer drones as the cost of implementation is quite high and there is also processing time needed. This will prevent GPS spoofing attacks but still would be vulnerable to jamming.

Even though arducopter did utilise IMU sensor fusion with GPS, we were still able to go undetected. On top of this, we can incorporate an additional sensor which will make the flight controller less reliant on GPS. For example, we can incorporate a camera as a visual odometry system which can track the drones movement in relation to the horizon line by comparing different frames provided by the camera. If the GPS data is indicating that the drone is moving but the frame hasn't changes and the horizon level is the same the system will flag it as an inconsistency and reject the data. Now with multiple cross-checks, this makes it much more difficult for an adversary to fool all the sensors.

If the drone has a data link, it can take use of external API's to validate its position. For example, by using googles Geolocation API the drone can send pings to multiple different cell towers, and the algorithm will be able to calculate its location to a very good accuracy. While this may induce

505 additional complexity and relies on a good connection, this will act as a
506 second layer of verification.

507 **6.3 Operational and policy strategies**

508 Many of the drones on the market already have a geo-fence databases such
509 as airports, military bases and prisons. Drone manufacturers need to have
510 a up to data no fly zone and a pre defined behaviour in case of a GPS loss.
511 This can reduce the impact caused by attacks. For example, if an adversary
512 lures the drone to an airport, a well defined no fly zone will get in the way
513 of travelling to such place.

514 Integration of spoofing detection in all UAV firmware as a standard will
515 raise the baseline security of drones. Making anti-spoofing features incor-
516 porated in all drones can narrow the window of opportunity for attackers.
517 This can include; standardised anomaly detection libraries or a certification
518 requirement that a drone is resilient to basic GPS spoof scenario.

519 Human operators can be put in the loop as a countermeasure. Organisa-
520 tion can make human operators an operational protocol and their responsi-
521 bility would be to monitor telemetry. If a drone starts to behave oddly or it
522 goes off course, the human operator can intervene and manually take control.
523 While this may not stop spoofing attacks, it provides situational awareness
524 so that the right steps can be taken.

525 Even though the UK government outlaws unauthorised signal transmis-
526 sion and jamming, not all countries have such strictly enforced laws. For
527 example, Ukraine is one of the few countries in the world that has legalised
528 the use of jammers. Funnily, the reason was so that they can be used to
529 stop students from using mobile phones during exams Phantom Technolo-
530 gies (2023).

531 Regulations alone cant stop a determined attacker, however, it can push
532 towards safer designs and ecosystems. By combining both technical defences
533 with operational policies, we can reduce the risk of GPS spoofing attacks
534 substantially. In the next section we will be going through adversaries can

benefit from GPS spoofing attacks.

6.4 Future works

In future work, I plan to explore this experiment on real hardware, potentially within an anechoic chamber to minimise signal interference and still maintaining within safety and ethical constraints. Additionally, I aim to refine the controller implementation to incorporate the burst technique we explored. However, given the existing delay, incorporating bursts could further complicate control.

I also intend to revisit the preliminary experiments to address the issues encountered with the script and the USRP B210. This will allow us to explore generating spoofed signals and simulating the communication based on the captured traces, which are crucial for achieving a realistic real-world attack conditions. Additionally, we will also explore the limitations imposed by rate of trace and the impact of environmental conditions.

The effectiveness of different countermeasure strategies can also be investigated in future work. The countermeasures can also be tested within the simulation environment we explored in this project.

7 Conclusion

GPS spoofing attacks against UAV's present a very serious and critical security threat when abused by malicious actors. Our projects success in hijacking a drone in a safe hardware free environment shows how adversaries can use GPS spoofing of consumer UAV's as a side-channel remote control method.

7.1 Adversarial threats and abuse scenarios

- Sabotaging delivery drones: An adversary can either steal a payload or simply just prevent a successful delivery by making the UAV go off course. For example, a competitor might spoof the drone off its pre

561 defined route, causing the drone to misdeliver or crash. This can have
562 reputational damages but also a possibility of getting your path/route
563 license revoked. This can cause a considerable amount of damage to
564 the customers trust of delivering parcels.

- 565 • Public safety: because GPS spoofing attacks allows us to lure the drone
566 to anywhere the attacker would like, an adversary can lure the drone
567 to a restricted or dangers airspace. For example, a malicious actor
568 can spoof a hobbyist drone to an airport such as Heathrow and cause
569 disruption or even accidents. We have seen in the past how a drones
570 sighting can halt all departures. In 2019, Heathrow airport was tem-
571 porarily stopped after a drone was reported to be seen BBC News
572 (2019).
- 573 • Attacks on critical infrastructure: drones that are being utilised for in-
574 specting infrastructure or by police can be spoofed to crash into critical
575 infrastructure such as power lines and bridges causing severe impact.
576 An adversary could also break formations of swarms by simply making
577 one drone collide with another without even requiring a large offset.

578 The application of GPS spoofing go beyond negative purposes, they also
579 have defensive and ethical applications.

- 580 • Safe-hijacking as threat mitigation: The concept of safe-hijacking can
581 be seen as an anti-drone defence method. Authorities can deliberately
582 spoof a rogue UAV and take remote control to be able to safely make it
583 land for further analysis. Simply jamming the rogue UAV can have the
584 opposite effect as the behaviour of the drone becomes unpredictable and
585 cause more damage. Noh et al. (2019) successful carried out a adaptive
586 GPS spoofing attack against commercial drones and it showed that
587 they can for the UAV to move in any direction with high accuracy in
588 real world conditions. Our simulation proved this possible even further.
589 With this accuracy, authorities can escort a rogue UAV out if harm's
590 way.

- Testing and training for resilience: This safe simulation-based approach provides drone developers and researchers a valuable tool to test their products. UAV manufacturers can evaluate their drones' robustness and fail-safe mechanisms under a GPS spoofing attack. For example, arducopter developers can adjust the EKF parameters or fine-tune the detection algorithm and run the spoofing scripts to see if the drone is still vulnerable to GPS spoofing attacks.

In conclusion, the project "Simulating GPS-spoofing of consumer UAVs as a side-channel remote control method" has illustrated the feasibility of GPS-based UAV hijacking, as well as to why it is important to guard against such attacks. While this research may not perfectly replicate the real-world conditions, it has demonstrated the drones' behaviour under fake GPS data and has highlighted bypass techniques for the Extended Kalman Filter algorithm.

In order for such an attack to be executed successfully in the real world, the attacker would require specific technical equipment, detailed knowledge of drones' flight controller and a malicious motive also including the factors of timing, rate of trace, signal precision and environmental conditions. Nevertheless, this project has effectively demonstrated the inherent vulnerability in UAV GPS navigation systems and has provided valuable insight and tools that can guide future developments in drone security and defensive measures.

7.2 Validation Framework for Enhancements

This simulation framework not only exploits the vulnerabilities in drone navigation under GPS spoofing but it can also be used to validate modifications to the navigation algorithm and sensor integration like the EKF. It enables developer to test incremental EKF adjustments and sensor fusion improvements, directly addressing the threats identified in previous sections. Researchers can benchmark different hardware upgrades like using better quality sensors and enhancements to the EKF algorithm to decide whether

algorithm can be developed that allows the use less accurate (legacy) sensors with improved efficacy when used as part of an enhanced EKF. This tool allows these theories and designs to be tested more rigorously.

References

- BBC News (2019). Gatwick drone attack possible inside job, say police. <https://www.bbc.co.uk/news/uk-46803713>, accessed: 2025-03-26.
- Gupta, S., Patnala, S. K., W, W., Agrawal, N., Tripathi, N., N, P. K. B. and Raman, B. (2023). Gps spoof and detect in ardupilot simulating uavs. In *Proceedings of the 21st OITS International Conference on Information Technology (OCIT 2023)*, IEEE.
- Humphreys, T. E., Ledvina, B. M., Psiaki, M. L., O’Hanlon, B. W. and Kintner, P. M. (2012). Assessing the spoofing threat: Development of a portable gps civilian spoofer. *IEEE Journal of Selected Topics in Signal Processing*, 5(3), pp. 574–581.
- Noh, J., Kwon, Y., Son, Y., Shin, H. and Kim, Y. (2019). Tractor beam: Safe-hijacking of consumer drones with adaptive gps spoofing. *ACM Transactions on Privacy and Security*, 22(2), pp. 12:1–12:26.
- Phantom Technologies (2023). Legality of Jammers. Accessed: 2025-03-26.